

1. Using the Big Data Cluster
 - Getting Access to the **bigdata** Cluster
1. Overview of Cloudera Services
2. Hadoop command line
 - Logging on to the Cluster via Hue
1. Using the HDFS file browser
2. Hosted datasets
3. Building an application
4. Wordcount in Spark
 - Read in a text file using the SparkContext
 1. Filter lines containing Spark
 - Map lines from String to Array[String]
 1. Scala Bonus!
 1. Let's split each line
 2. Counting words per line
 3. Counting words per line
 - Counting total "Spark" occurrences with **flatMap**
 - Counting occurrences of each word
 1. Reduce

1 Using the Big Data Cluster

ACCRE currently has a development environment set up for Vanderbilt and VUMC researchers to access, and plan to build out a production environment over the next 2-3 years. The development environment, (aka **bigdata**) is running the Cloudera implementation of the Hadoop ecosystem.

2 Getting Access to the bigdata Cluster

The bigdata cluster is available for use by the Vanderbilt community. Users should contact ACCRE to get access to the cluster.

3 Overview of Cloudera Services

Service	Description
YARN	Yet Another Resource Negotiator
Oozie	Web app for scheduling Hadoop jobs
MapReduce 2	MapReduce jobs running on top of YARN
Hue	User interface for constructing Jobs

Service	Description
Spark	MapReduce-like + cacheing
Hive	ETL transformations expressed as SQL
Impala	Interactive SQL
HBase	Random, realtime read/write access to distributed big data store
Pig	High-level language for expressing data analysis programs
Solr	Text search engine supporting free form queries

4 Hadoop command line

To interact with the HDFS filesystem, use the `hadoop fs` command:

```
[arnoldjr@abd740 ~]$ hadoop fs -ls /
Found 5 items
drwxr-xr-x - hdfs supergroup          0 2017-04-19 13:47 /data
drwxr-xr-x - hbase hbase              0 2017-04-02 21:09 /hbase
drwxrwxr-x - solr solr                0 2017-02-24 17:20 /solr
drwxrwxrwx - hdfs supergroup          0 2017-05-06 00:26 /tmp
drwxr-xr-x - hdfs supergroup          0 2017-02-17 12:14 /user
[arnoldjr@abd740 ~]$ hadoop fs -ls /data
Found 9 items
-rw-r--r--  3 hdfs supergroup    3359 2017-02-14 09:57 /data/Spark_README.md
drwxr-xr-x - hdfs supergroup      0 2017-03-06 16:25 /data/babs
drwxr-xr-x - hdfs supergroup      0 2017-03-06 11:52 /data/capitalbikeshare-data
drwxr-xr-x - hdfs supergroup      0 2017-03-06 12:10 /data/citibike-tripdata
drwxr-xr-x - hdfs supergroup      0 2017-02-14 21:10 /data/google-ngrams
-rw-r--r--  3 hdfs supergroup 274188932 2017-04-19 13:47 /data/hadoop-2.5.0-cdh5.2.0.tar.gz
drwxr-xr-x - hdfs supergroup      0 2017-01-18 19:06 /data/nyc-tlc
drwxr-xr-x - hdfs supergroup      0 2016-12-21 15:14 /data/stack-archives
```

5 Logging on to the Cluster via Hue

Once approved, users will be able to connect to `bigdata.accre.vanderbilt.edu` via `ssh`, but Cloudera Manager provides a WebUI to interact with the cluster called Hue. To access Hue, simply to `bigdata.accre.vanderbilt.edu:8888` in your web browser and enter your credentials.

6 Using the HDFS file browser

If you've used the web UIs for Dropbox, Google Drive, etc., then this step is a piece of cake. The File Browser is accessed from the dog-eared-piece-of-paper

icon near the top right of the screen. In the file browser, you're able to navigate the directory structure of HDFS and even view the contents of text files.

- When a new user logs into Hue, Hue creates an HDFS directory for that user at `/user/<vunetid>` which becomes that user's home directory.
- *Note that, by default, logging in to Hue creates a new user's home directory with read and execute permissions enabled for the world!*
- Files can be uploaded to your directories using the drag-and-drop mechanism; however, the file size for transferring through the WebUI is capped at around 50GB, so other tools like `scp` or `rsync` are necessary for moving large files onto the cluster.

7 Hosted datasets

In addition to your own data, ACCRE hosts some publicly available datasets at `/data/`:

Directory	Description
babs	Bay Area bikeshare data
capitalbikeshare data	DC area bikeshare data
citibike-tripdata	NYC bikeshare data
google-ngrams	n-grams collected from Google Books
nyc-tlc	NYC taxi trip data
stack-archives	historic posts from StackOverflow, et al.

If you know of other datasets that may appeal to the Vanderbilt community at large, just let us know!

8 Building an application

Hue uses Oozie to compose workflows on the cluster, and to access it, you'll need to follow the tabs `Workflows -> Editors -> Workflows`.

From here, click the **+ Create** button, and you'll arrive at the workflow composer screen. You can drag and drop an application into your workflow, for instance a Spark job. Here you can specify the jar file (which, conveniently, you can generate from our GitHub repo) and specify options and inputs.

If you want to interactively select your input and output files each time you execute the job, you can use the special keywords `${input}` and `${output}`, which is a nice feature for generalizing your workflows.

9 Wordcount in Spark

Here, we present an example Spark application. This content is adapted slightly from the Spark getting started guide. Users can execute the same commands in the Spark REPL, which is launched by running in bash

```
$SPARK_HOME/bin/spark-shell
```

where the environment variable `$SPARK_HOME` points to an installation of Spark. This can be a local Spark installation, or on `bigdata`, users can type `spark-shell` for Spark 1.6.2 or `spark2-shell` for Spark 2.0.0.

9.1 Read in a text file using the SparkContext

- The SparkContext `sc` is the entry point for Spark's data structures (provided automatically in the Spark REPL).
- The value `textFile` becomes an RDD (Resilient Distributed Dataset)

```
val textFile = sc.textFile("spark_read_me.txt")
```

- The RDD cannot be viewed directly in the REPL (in practice it will be distributed across many nodes!!). Thus, in order to view all the data, we have to gather the data at a single node using `collect`. A summary of RDD functions can be found [here](#).

```
(textFile collect) foreach println
```

```
# Apache Spark
```

Spark is a fast and general cluster computing system for Big Data. It provides high-level APIs in Scala, Java, Python, and R, and an optimized engine that supports general computation graphs for data analysis. It also supports a rich set of higher-level tools including Spark SQL for SQL and DataFrames, MLlib for machine learning, GraphX for graph processing, and Spark Streaming for stream processing.

```
<http://spark.apache.org/>
```

9.2 Filter lines containing Spark

- The RDD (and Scala collections) support filtering

```
val linesWithSpark = textFile.filter(line => line.contains("Spark"))  
(linesWithSpark collect) foreach println
```

```
# Apache Spark
Spark is a fast and general cluster computing system for Big Data. It provides
rich set of higher-level tools including Spark SQL for SQL and DataFrames,
and Spark Streaming for stream processing.
You can find the latest Spark documentation, including a programming
## Building Spark
Spark is built using [Apache Maven](http://maven.apache.org/).
To build Spark and its example programs, run:
["Building Spark"](http://spark.apache.org/docs/latest/building-spark.html).
The easiest way to start using Spark is through the Scala shell:
Spark also comes with several sample programs in the `examples` directory.
    ./bin/run-example SparkPi
    MASTER=spark://host:7077 ./bin/run-example SparkPi
Testing first requires [building Spark](#building-spark). Once Spark is built, tests
Spark uses the Hadoop core library to talk to HDFS and other Hadoop-supported
Hadoop, you must build Spark against the same version that your cluster runs.
in the online documentation for an overview on how to configure Spark.
```

9.3 Map lines from String to Array[String]

The RDD (and Scala collections) support mapping. For example:

```
val foo = "a line with Spark"
foo.split(" ")
Array(a, line, with, Spark)
```

—

9.3.1 Scala Bonus!

The dot operator can be omitted in Scala, so that splitting operation can be written as:

```
foo split (" ")
Array(a, line, with, Spark)
```

9.4 Let's split each line

```
val arraysWithSpark = linesWithSpark.map(line => line split(" "))
val wordsPerLine = arraysWithSpark map (a => a.size)
(wordsPerLine collect) foreach println
```

```
3
14
12
6
10
3
6
.
.
.
```

9.5 Counting words per line

- Spark and Scala allow for chaining operations together. Thus, we can just write:

```
val wordsPerLine = linesWithSpark.map(line => line split(" ")).map(a => a.size)
(wordsPerLine collect) foreach println
```

```
3
14
12
6
10
3
6
.
.
.
```

9.6 Counting words per line

- Equivalently:

```
val wordsPerLine = linesWithSpark map(_ split " ") map(_ size)
(wordsPerLine collect) foreach println
```

```
3
14
12
6
10
3
6
.
```

.
.

Here, Scala is smart enough to understand that the underscore implies that the map function is taking a single argument, thus avoiding the need for the `foo => foo split(" ")` pattern.

Scala Bonus! this pattern defines an anonymous function.

9.7 Counting total “Spark” occurrences with flatMap

If we wanted to count total occurrence of *Spark* we could filter our `Arrays` from the previous step to keep only those words that match *Spark*. Since we don’t actually care about which line contains occurrences of *Spark* but rather how many *Sparks* are in our entire document, we can consider all the individual words at once. Spark provides a mechanism to do this called `flatMap`.

```
val allSparkWords = linesWithSpark flatMap (line => line split " ")
(allSparkWords collect) foreach println
```

```
#
Apache
Spark
Spark
is
a
fast
and
general
cluster
computing
system
for
Big
.
.
.
```

To get the Spark occurrences, simply filter:

```
val allSparkWords = linesWithSpark flatMap (_ split "\\s+") filter (_ == "Spark")
(allSparkWords collect) foreach println
```

```
Spark
Spark
Spark
Spark
Spark
Spark
Spark
```

Spark

Spark

.
.
.

Scala Bonus! "\\s+" is a regex that matches one or more consecutive whitespace characters.

9.8 Counting occurrences of each word

- The classic MapReduce example (as popularized by Hadoop)
- Split each line into words
- Map each word into a word value *pair*, or *Tuple*, e.g. (*the*, 100)
- First element is the *key* which serves as identifier
- Second element is the *value* which signifies that each word has occurred one time.

Since each word is not unique, we need to group them together and count the occurrences per group. When we perform an *action* on an RDD, all the pairs with identical keys are sent to the same node and then we can aggregate these together. This is precisely what the `reduceByKey` function does:

```
val wordFrequencies = textFile flatMap (_ split "\\s+") map (word => (word, 1)) reduceByKey  
(wordFrequencies take 20) foreach println  
  
(package,1)  
(this,1)  
(Version,"(http://spark.apache.org/docs/latest/building-spark.html#specifying-the-hadoop-ver  
(Because,1)  
(Python,2)  
(cluster.,1)  
(its,1)  
.  
.  
.
```

9.9 Reduce

The argument to `reduceByKey` is a function describing how to combine values (which must have the same type, otherwise see `aggregateByKey`). To print the output, we've used the `take` function to take the first 20 results. We can also use the aptly named `takeOrdered` function:

```
(wordFrequencies.takeOrdered(50)(Ordering[Int].reverse.on(_._2))) foreach println  
(,43)  
(the,21)
```


(to,14)
(Spark,13)
(for,11)
(and,10)
(a,8)
(##,8)
(run,7)
(is,6)
(can,6)
(on,5)
(in,5)
(of,5)
(also,4)
.
.
.