


Introduction to Using Spark on the Big Data Cluster

Getting Help

- Submit a ticket to [the Helpdesk](#)
- Join the accre-forum Slack team!
 - <https://accre-forum.slack.com/signup>
 - Use your Vanderbilt email address to register
 - Join the #bigdata channel to communicate and collaborate
- Check out our GitHub organization at github.com/bigdata-vandy
- Check out our blog at bigdata-vandy.github.io
- Schedule a meeting for your research group

The Hadoop Ecosystem



How is Big Data different from HPC?

Traditional HPC

Move the data to the code

Centralized storage via GPFS

Perfomant languages (C/C++, Fortran)

Imperative programming

Big Data

Move the code to the data

Distributed storage via HDFS

Abstract languages (Java, Scala, Python)

Functional programming

How is Big Data different from HPC?

Traditional HPC

Move the data to the code

Centralized storage via GPFS

Performant languages (C/C++, Fortran)

Imperative programming

Big Data

Move the code to the data

Distributed storage via HDFS

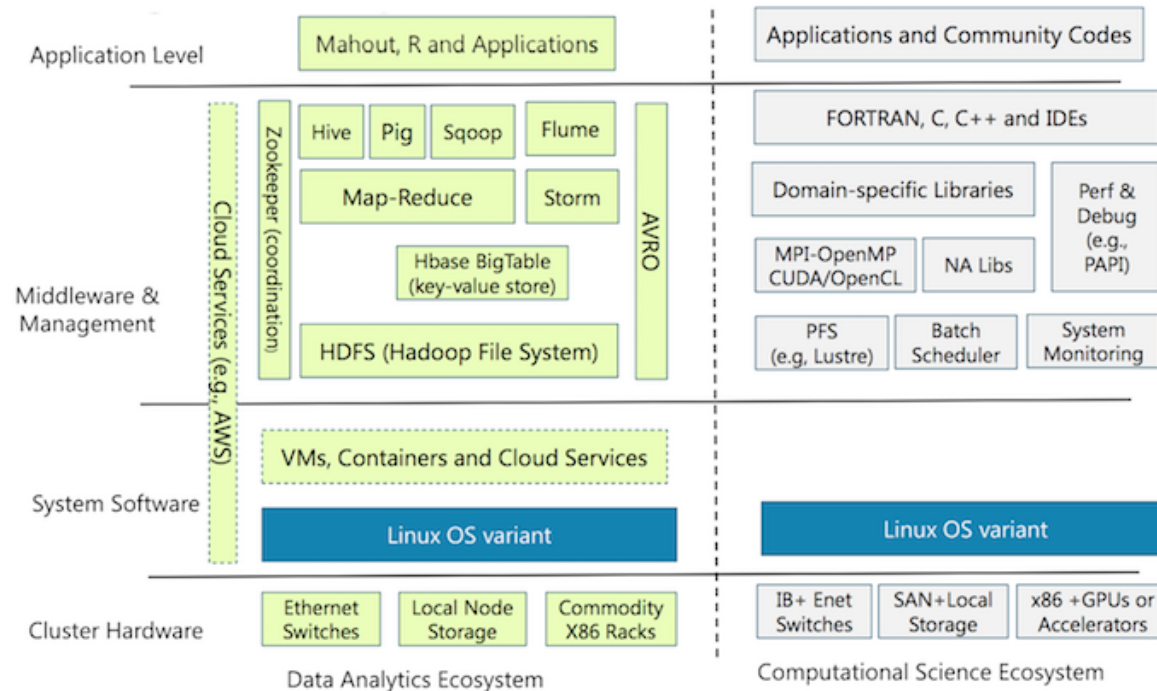
Abstract languages (Java, Scala, Python)

Functional programming

Candidates for Big Data solutions:

- Lots of data
- Embarrassingly parallel data processing
- Machine learning
- Network analysis
- Processing/analysis of streaming data

Divergent Ecosystems



Dan Reed, "Clouds, Big Data, and the Future of Computing," CASC Fall 2015 Meeting, Washington, D.C.

BigData Test Cluster - Available Now!

Initial deployment for testing, prototyping, and benchmarking only, using recycled hardware

- 3 management nodes
 - 8 CPU cores, 92 GB RAM per node
- 6 data nodes
 - 8 TB, 8 CPU cores, 92 GB RAM per node

BigData Test Cluster - Available Now!

Initial deployment for testing, prototyping, and benchmarking only, using recycled hardware

- 3 management nodes
 - 8 CPU cores, 92 GB RAM per node
- 6 data nodes
 - 8 TB, 8 CPU cores, 92 GB RAM per node

BigData Production Environment - Available Fall 2017

Newly purchased hardware!!

- 4 management nodes
 - 16 CPU cores, 128 GB RAM per node
- 9 data nodes
 - 80 TB, 32 CPU cores, 512 GB RAM per node

BigData Cluster Management with Cloudera

Cloudera Manager is used to deploy, configure, and manage the bigdata cluster

BigData Cluster Management with Cloudera

Cloudera Manager is used to deploy, configure, and manage the bigdata cluster

Service	Description
HDFS	Hadoop Distributed File System - replicated, partitioned data
YARN	Yet Another Resource Negotiator
MapReduce 2	MapReduce jobs running on top of YARN
Spark	MapReduce-like + cacheing
Oozie	Web app for scheduling Hadoop jobs
Hue	User interface for constructing Jobs
Hive	ETL transformations expressed as SQL
Impala	Interactive SQL
HBase	Random, realtime read/write access to distributed big data store
Pig	High-level language for expressing data analysis programs
Solr	Text search engine supporting free form queries

Technologies for this Presentation

Service Function

HDFS Store data in a distributed/replicated manner

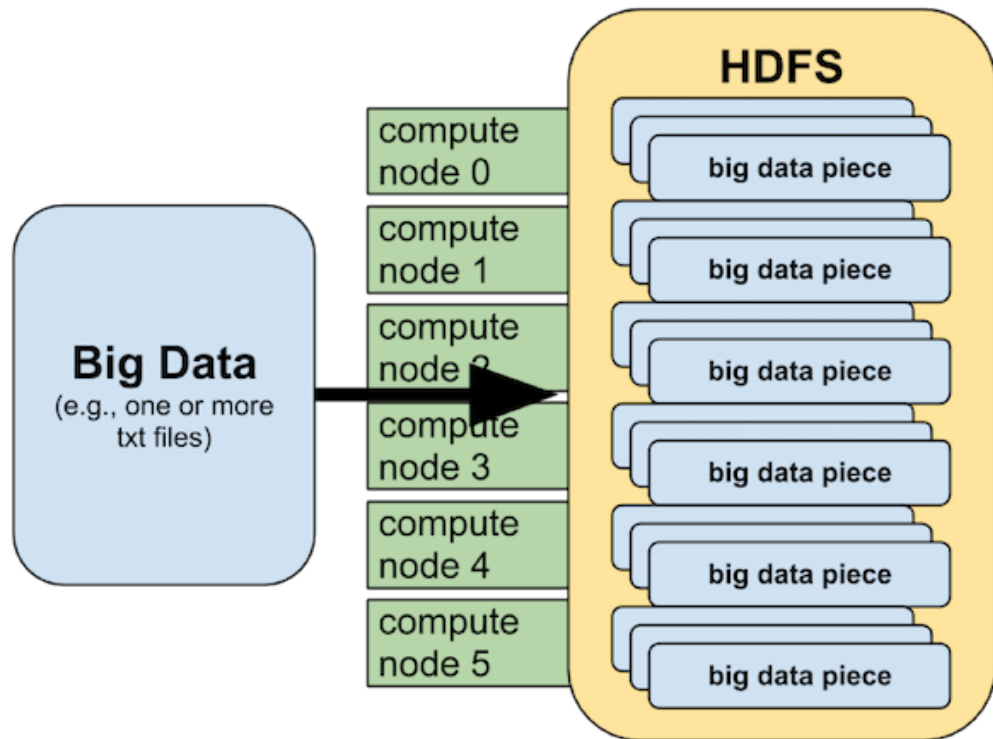
YARN Schedule and run jobs, acquiring resources as necessary

Spark Perform map-reduce with data persistence in memory + much, much more

What is HDFS?

- Distributed file system
- Designed to run on commodity hardware
- Highly fault-tolerant (through replication)
- Suitable for applications that have large data sets

HDFS Magic



Interacting with HDFS

The HDFS filesystem uses Unix-like commands for common operations, prefixed with the `hadoop fs` command, e.g.:

```
$ hadoop fs -ls /  
Found 5 items  
drwxr-xr-x   - hdfs  supergroup          0 2017-04-19 13:47 /data  
drwxr-xr-x   - hbase hbase              0 2017-04-02 21:09 /hbase  
drwxrwxr-x   - solr  solr              0 2017-02-24 17:20 /solr  
drwxrwxrwx   - hdfs  supergroup          0 2017-05-06 00:26 /tmp  
drwxr-xr-x   - hdfs  supergroup          0 2017-02-17 12:14 /user
```

Interacting with HDFS

The HDFS filesystem uses Unix-like commands for common operations, prefixed with the `hadoop fs` command, e.g.:

```
$ hadoop fs -ls /
Found 5 items
drwxr-xr-x   - hdfs  supergroup          0 2017-04-19 13:47 /data
drwxr-xr-x   - hbase hbase              0 2017-04-02 21:09 /hbase
drwxrwxr-x   - solr  solr              0 2017-02-24 17:20 /solr
drwxrwxrwx   - hdfs  supergroup          0 2017-05-06 00:26 /tmp
drwxr-xr-x   - hdfs  supergroup          0 2017-02-17 12:14 /user

$ hadoop fs -ls /data
Found 9 items
-rw-r--r--   3 hdfs  supergroup    3359 2017-02-14 09:57 /data/Spark_README.md
drwxr-xr-x   - hdfs  supergroup          0 2017-03-06 16:25 /data/babs
drwxr-xr-x   - hdfs  supergroup          0 2017-03-06 11:52 /data/capitalbikeshare-data
drwxr-xr-x   - hdfs  supergroup          0 2017-03-06 12:10 /data/citibike-tripdata
drwxr-xr-x   - hdfs  supergroup          0 2017-02-14 21:10 /data/google-ngrams
-rw-r--r--   3 hdfs  supergroup 274188932 2017-04-19 13:47 /data/hadoop-2.5.0-cdh5.2.0.tar.gz
drwxr-xr-x   - hdfs  supergroup          0 2017-01-18 19:06 /data/nyc-tlc
drwxr-xr-x   - hdfs  supergroup          0 2016-12-21 15:14 /data/stack-archives
```

Interacting with HDFS

The HDFS filesystem uses Unix-like commands for common operations, prefixed with the `hadoop fs` command, e.g.:

```
$ hadoop fs -ls /
Found 5 items
drwxr-xr-x   - hdfs  supergroup          0 2017-04-19 13:47 /data
drwxr-xr-x   - hbase hbase              0 2017-04-02 21:09 /hbase
drwxrwxr-x   - solr  solr              0 2017-02-24 17:20 /solr
drwxrwxrwx   - hdfs  supergroup          0 2017-05-06 00:26 /tmp
drwxr-xr-x   - hdfs  supergroup          0 2017-02-17 12:14 /user

$ hadoop fs -ls /data
Found 9 items
-rw-r--r--   3 hdfs  supergroup    3359 2017-02-14 09:57 /data/Spark_README.md
drwxr-xr-x   - hdfs  supergroup          0 2017-03-06 16:25 /data/babs
drwxr-xr-x   - hdfs  supergroup          0 2017-03-06 11:52 /data/capitalbikeshare-data
drwxr-xr-x   - hdfs  supergroup          0 2017-03-06 12:10 /data/citibike-tripdata
drwxr-xr-x   - hdfs  supergroup          0 2017-02-14 21:10 /data/google-ngrams
-rw-r--r--   3 hdfs  supergroup 274188932 2017-04-19 13:47 /data/hadoop-2.5.0-cdh5.2.0.tar.gz
drwxr-xr-x   - hdfs  supergroup          0 2017-01-18 19:06 /data/nyc-tlc
drwxr-xr-x   - hdfs  supergroup          0 2016-12-21 15:14 /data/stack-archives
```

This is not the same /data as on GPFS

HDFS

Common arguments for the [hadoop fs command](#):

Argument	Description
cat	concatenate file contents to stdout
copyFromLocal	copy from local filesystem to HDFS
copyToLocal	copy from HDFS to local disk
du	calculate disk usage
moveFromLocal	copy to HDFS and remove local copy
moveToLocal	copy to local drive and remove HDFS copy
rm	remove files and folders from HDFS

HDFS

Common arguments for the [hadoop fs command](#):

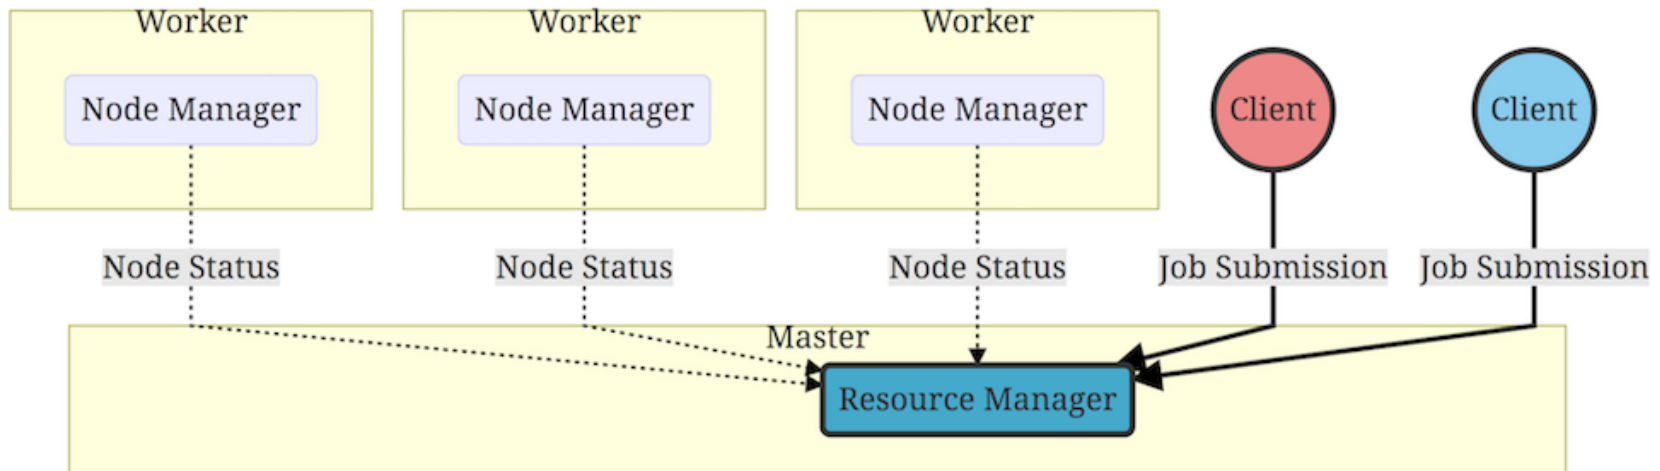
Argument	Description
cat	concatenate file contents to stdout
copyFromLocal	copy from local filesystem to HDFS
copyToLocal	copy from HDFS to local disk
du	calculate disk usage
moveFromLocal	copy to HDFS and remove local copy
moveToLocal	copy to local drive and remove HDFS copy
rm	remove files and folders from HDFS

Command outputs can be piped, just like normal Unix commands, for example

```
hadoop fs -cat /user/fido/my-file | head -100
```

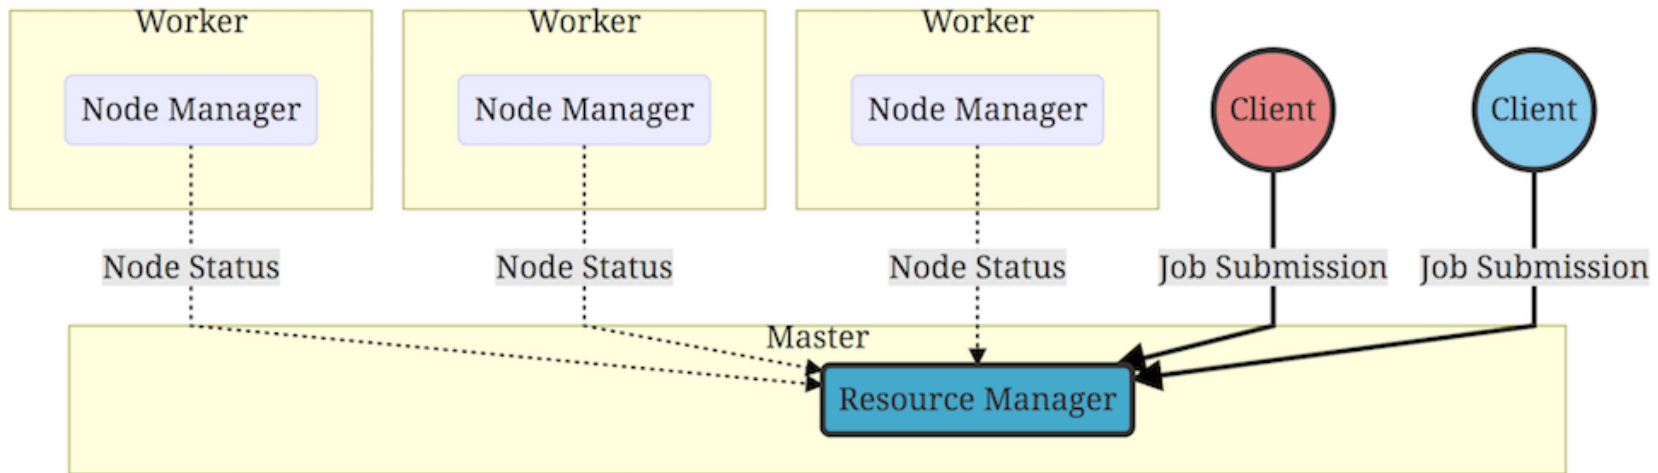
Apache Hadoop YARN

- Yet Another Resource Negotiator
- Functionalities:
 - resource management
 - job scheduling/monitoring



Apache Hadoop YARN

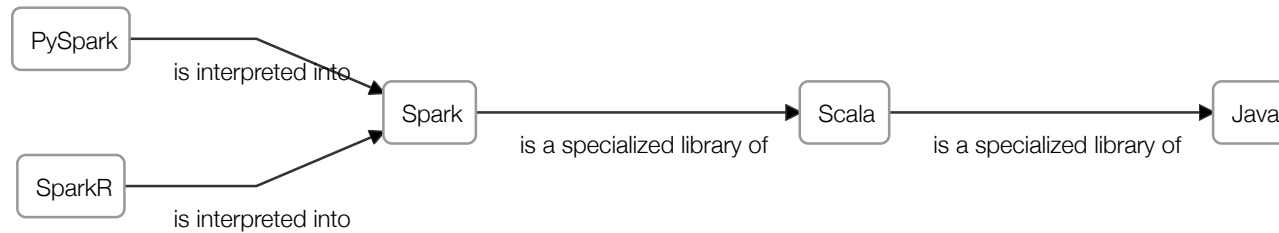
- Yet Another Resource Negotiator
- Functionalities:
 - resource management
 - job scheduling/monitoring



- YARN effectively fills the same role as SLURM
- Resource allocation is essentially transparent to the user

What is Spark?

Apache Spark is a fast and general-purpose cluster computing system. It provides high-level APIs in Java, Scala, Python and R, and an optimized engine that supports general execution graphs.



Spark can run in:

- Standalone mode, e.g. [Spark on GPFS](#)
- on YARN
- on Mesos

How to execute Spark code

- Interactive jobs: Spark REPL, a command line tool to "Read Evaluate Print Loop" Spark/Scala code
 - In general: `$SPARK_HOME/bin/spark-shell`
 - On the cluster:
 - `spark-shell` or `pyspark` (v. 1.6.0)
 - `spark2-shell` or `pyspark2` (v. 2.0.0)

How to execute Spark code

- Interactive jobs: Spark REPL, a command line tool to "Read Evaluate Print Loop" Spark/Scala code
 - In general: `$SPARK_HOME/bin/spark-shell`
 - On the cluster:
 - `spark-shell` or `pyspark` (v. 1.6.0)
 - `spark2-shell` or `pyspark2` (v. 2.0.0)
- Batch jobs:
 - In general: `$SPARK_HOME/bin/spark-submit`
 - On the cluster:
 - `spark-submit <jar-file>` or `spark-submit <python-file>` (v. 1.6.0)
 - `spark2-submit <jar-file>` or `spark2-submit <python-file>` (v. 2.0.0)

How to execute Spark code

- Interactive jobs: Spark REPL, a command line tool to "Read Evaluate Print Loop" Spark/Scala code
 - In general: `$SPARK_HOME/bin/spark-shell`
 - On the cluster:
 - `spark-shell` or `pyspark` (v. 1.6.0)
 - `spark2-shell` or `pyspark2` (v. 2.0.0)
- Batch jobs:
 - In general: `$SPARK_HOME/bin/spark-submit`
 - On the cluster:
 - `spark-submit <jar-file>` or `spark-submit <python-file>` (v. 1.6.0)
 - `spark2-submit <jar-file>` or `spark2-submit <python-file>` (v. 2.0.0)


On the bigdata cluster, `spark-submit` calls on YARN to automatically distributes jobs.

Wordcount in Spark

- Wordcount is the "Hello World" application for Hadoop
- Count the occurrences of all words in a text file
- Use the Scala REPL

Wordcount in Spark

- Wordcount is the "Hello World" application for Hadoop
- Count the occurrences of all words in a text file
- Use the Scala REPL
- This content is adapted slightly from the [Spark getting started guide](#).

 Overview Programming Guides ▾ API Docs ▾ Deploying ▾ More ▾

Quick Start

- [Interactive Analysis with the Spark Shell](#)
 - [Basics](#)
 - [More on RDD Operations](#)
 - [Caching](#)
- [Self-Contained Applications](#)
- [Where to Go from Here](#)

This tutorial provides a quick introduction to using Spark. We will first introduce the API through Spark's interactive shell (in Python or Scala), then show how to write applications in Java, Scala, and Python. See the [programming guide](#) for a more complete reference.

Read in a text file using the SparkContext

- The SparkContext class `sc` is the entry point for the Spark API, and is already created in the REPL

```
scala> val linesRDD = sc.textFile("spark_read_me.txt")  
linesRDD: org.apache.spark.rdd.RDD[String] = spark_read_me.txt MapPartitionsRDD[1] at textFile at <console>:24
```

Read in a text file using the SparkContext

- The SparkContext class `sc` is the entry point for the Spark API, and is already created in the REPL

```
scala> val linesRDD = sc.textFile("spark_read_me.txt")
linesRDD: org.apache.spark.rdd.RDD[String] = spark_read_me.txt MapPartitionsRDD[1] at textFile at <console>:24
```

- The value `linesRDD` is an instance of an RDD (Resilient Distributed Dataset)

Read in a text file using the SparkContext

- The SparkContext class `sc` is the entry point for the Spark API, and is already created in the REPL

```
scala> val linesRDD = sc.textFile("spark_read_me.txt")
linesRDD: org.apache.spark.rdd.RDD[String] = spark_read_me.txt MapPartitionsRDD[1] at textFile at <console>:24
```

- The value `linesRDD` is an instance of an RDD (Resilient Distributed Dataset)
- RDDs are essential to Spark
 - a fault-tolerant collection of elements that can be operated on in parallel
 - parallelism happens *automagically*!

View the contents of the linesRDD RDD

- Since RDDs are distributed across nodes, they cannot be viewed directly in the REPL
- To view an RDD, gather all the data at a single node using collect.

```
scala> linesRDD.collect.foreach(println)
# Apache Spark
```

Spark is a fast and general cluster computing system **for Big Data**. It provides high-level **APIs** in **Scala**, **Java**, **Python**, and **R**, and an optimized engine that supports general computation graphs **for** data analysis. It also supports a rich set of higher-level tools including **Spark SQL for SQL** and **DataFrames**, **MLlib for** machine learning, **GraphX for** graph processing, and **Spark Streaming for** stream processing.

<<http://spark.apache.org/>>

.
.br/>.

Map lines from String to Array[String]

The RDD (and Scala collections) support mapping. For example, the pure Scala expression:

```
scala> val foo: String = "a line with Spark"
```

Can be mapped into an array of strings delimited by whitespace

```
scala> foo.split(" ")  
res0: Array[String] = Array(a, line, with, Spark)
```

Counting words per line

The same splitting operation can be applied to each element of the RDD[String]

```
scala> val words: RDD[Array[String]] = linesRDD.map(line => line.split(" "))  
words: org.apache.spark.rdd.RDD[Array[String]] = MapPartitionsRDD[3] at map at <console>:27
```


Counting words per line

The same splitting operation can be applied to each element of the RDD[String]

```
scala> val words: RDD[Array[String]] = linesRDD.map(line => line.split(" "))
words: org.apache.spark.rdd.RDD[Array[String]] = MapPartitionsRDD[3] at map at <console>:27
```

Taking the size of each Array[String] element of the RDD will yield the total words per each line

```
scala> val wordsPerLine: RDD[Int] = words.map(a => a.size)
wordsPerLine: org.apache.spark.rdd.RDD[Int] = MapPartitionsRDD[5] at map at <console>:29
```

```
scala> wordsPerLine.take(5).foreach(println)
```

3

1

14

13

11

Total words

To get the total words for the document, we can simply sum up the wordsPerLine RDD:

```
scala> val total: Int = wordsPerLine.reduce((a, b) => a + b)
total: Int = 507
```

Total words

To get the total words for the document, we can simply sum up the wordsPerLine RDD:

```
scala> val total: Int = wordsPerLine.reduce((a, b) => a + b)
total: Int = 507
```

- The function passed to `reduce` specifies how to combine two `Int` elements.
- The order of combining the RDD elements must not matter

Total words

To get the total words for the document, we can simply sum up the wordsPerLine RDD:

```
scala> val total: Int = wordsPerLine.reduce((a, b) => a + b)
total: Int = 507
```

- The function passed to `reduce` specifies how to combine two `Int` elements.
- The order of combining the RDD elements must not matter

Notice that `total` is no longer an RDD but a regular Scala `Int`.

The MapReduce Paradigm

- Map: transform the data
 - map
 - filter
 - flatMap
 - join
 - union
 - intersection
 - et al.
- Reduce: aggregate the data
 - reduce
 - count
 - sample
 - et al.

WordCount: Count the occurrences of each word in a document

- Pseudocode:
 - Map each line into an array of words
 - Map each word into a word-value *pair*, or Tuple, e.g. (the, 101)
 - First element is the *key* which serves as identifier
 - Second element is the *value* which, in this case, signifies that each word has occurred one time.
 - Group all word-value pairs by word, and sum up the values for each word

WordCount: Count the occurrences of each word in a document

- Pseudocode:
 - Map each line into an array of words
 - Map each word into a word-value *pair*, or Tuple, e.g. (the, 101)
 - First element is the *key* which serves as identifier
 - Second element is the *value* which, in this case, signifies that each word has occurred one time.
 - Group all word-value pairs by word, and sum up the values for each word
- Key-Value pairs: mechanism for allowing `groupBy` operations
 - All the RDD elements of a given key are guaranteed to on the same worker during aggregations
 - `reduceByKey` combines values of a given key:

WordCount: Count the occurrences of each word in a document

- Pseudocode:
 - Map each line into an array of words
 - Map each word into a word-value *pair*, or Tuple, e.g. (the, 101)
 - First element is the *key* which serves as identifier
 - Second element is the *value* which, in this case, signifies that each word has occurred one time.
 - Group all word-value pairs by word, and sum up the values for each word
- Key-Value pairs: mechanism for allowing `groupBy` operations
 - All the RDD elements of a given key are guaranteed to on the same worker during aggregations
 - `reduceByKey` combines values of a given key:

```
scala> val wordFrequencies = linesRDD.flatMap( line => line.split(" ").map(word => (word, 1)) ).  
    | reduceByKey((a, b) => a + b)  
wordFrequencies: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[7] at reduceByKey at <console>:28
```

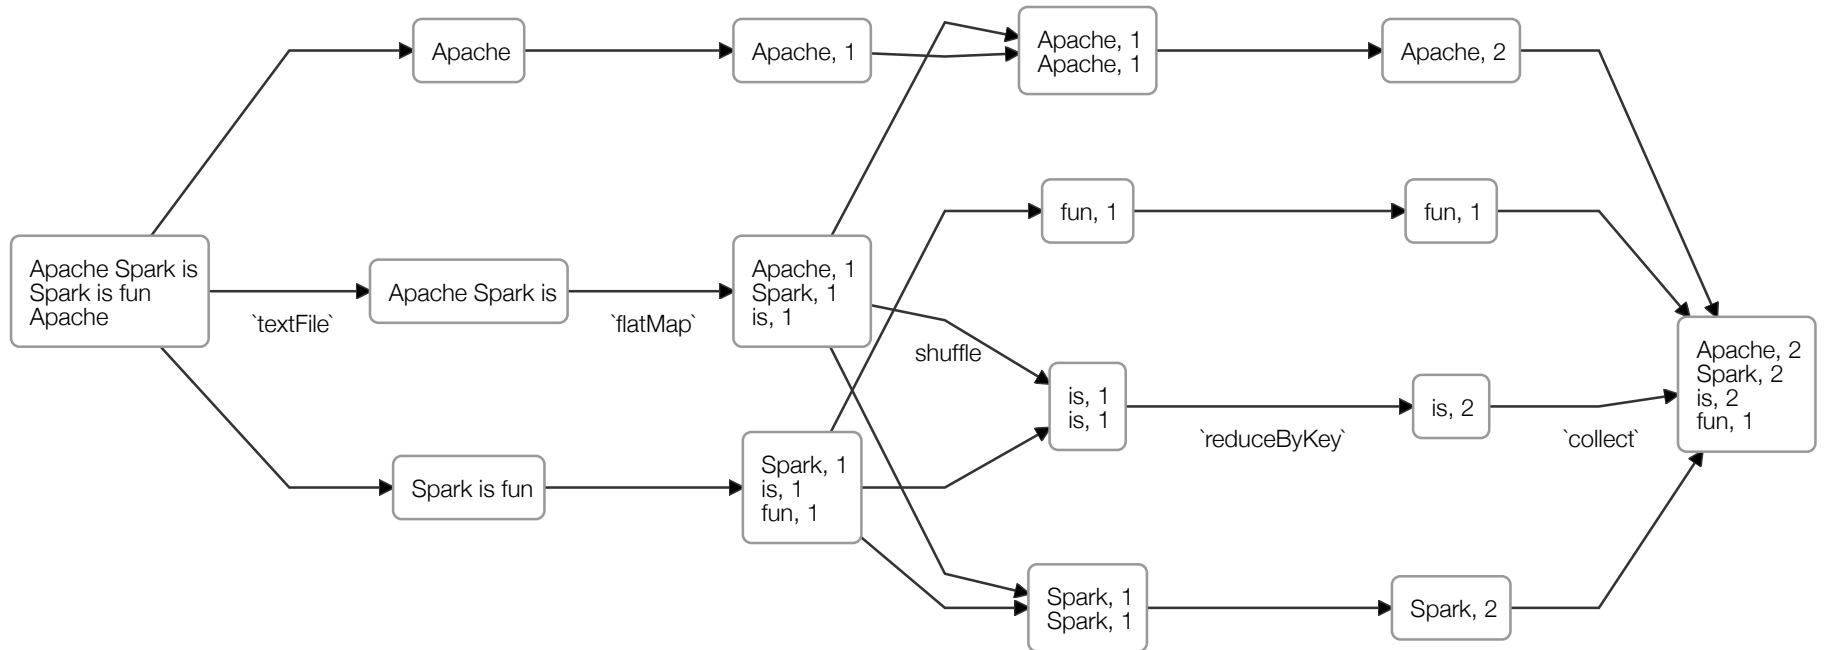
--

Note that `wordFrequencies` is now a `ShuffledRDD`

WordCount

```
scala> wordFrequencies.takeOrdered(20).foreach(println)
(,67)
("local",1)
("local[N]",1)
("yarn",1)
(#,1)
(##,8)
((You,1)
(-DskipTests,1)
(./bin/pyspark,1)
(./bin/run-example,2)
(./bin/spark-shell,1)
(./dev/run-tests,1)
(1000).count(),1)
(1000:,2)
(<class>,1)
(<http://spark.apache.org/>,1)
(>>>,1)
(A,1)
(APIs,1)
(About,1)
```

Under the Hood of Map-Reduce



Spark higher-level tools

Library	Function
Spark SQL	SQL and structured data processing
MLlib	machine learning
GraphX	graph processing
Spark Streaming	real-time analysis of data streams

Spark DataFrames

- Analogous to data frames in R and Pandas
- Can be automatically created from CSV, Parquet, JSON lines files

```
val df: DataFrame = spark.read
  .option("header", true)
  .option("treatEmptyValuesAsNulls", true)
  .option("inferSchema", true)
  .csv("/data/nyc-tlc/yellow_tripdata_2016-04.csv")
```

```
df.show()
```

vendor_id	pickup_datetime	dropoff_datetime	passenger_count	trip_distance	pickup_longitude
CMT	2013-01-01 15:11:48	2013-01-01 15:18:10	4	1.0	-73.978165
CMT	2013-01-06 00:18:35	2013-01-06 00:22:54	1	1.5	-74.00668
CMT	2013-01-05 18:49:41	2013-01-05 18:54:23	1	1.1	-74.004711
CMT	2013-01-07 23:54:15	2013-01-07 23:58:20	2	0.7	-73.9746
CMT	2013-01-07 23:25:03	2013-01-07 23:34:24	1	2.1	-73.976252
CMT	2013-01-07 15:27:48	2013-01-07 15:38:37	1	1.7	-73.966743
CMT	2013-01-08 11:01:15	2013-01-08 11:08:14	1	0.8	-73.995801
CMT	2013-01-07 12:39:18	2013-01-07 13:10:56	3	10.7	-73.989937
CMT	2013-01-07 18:15:47	2013-01-07 18:20:47	1	0.8	-73.980071
CMT	2013-01-07 15:33:28	2013-01-07 15:49:26	2	2.5	-73.977937
CMT	2013-01-08 13:11:52	2013-01-08 13:19:50	1	1.3	-73.982456
CMT	2013-01-08 09:50:05	2013-01-08 10:02:54	1	0.7	-73.995561

SQL queries on Spark DataFrames

```
scala> val coords = df.select("pickup_longitude", "pickup_latitude")
coords: org.apache.spark.sql.DataFrame = [pickup_longitude: string, pickup_latitude: string]
scala> coords.show(5)
+-----+-----+
| pickup_longitude| pickup_latitude|
+-----+-----+
|-73.990211486816406|40.750968933105469|
|-73.935188293457031| 40.80072021484375|
|-73.963752746582031|40.767936706542969|
|-73.997177124023438|40.742168426513672|
|-74.006843566894531|40.730266571044922|
+-----+-----+
only showing top 5 rows
```

SQL queries on Spark DataFrames

```
scala> val coords = df.select("pickup_longitude", "pickup_latitude")
coords: org.apache.spark.sql.DataFrame = [pickup_longitude: string, pickup_latitude: string]
scala> coords.show(5)
```

```
+-----+-----+
| pickup_longitude| pickup_latitude|
+-----+-----+
|-73.990211486816406|40.750968933105469|
|-73.935188293457031| 40.80072021484375|
|-73.963752746582031|40.767936706542969|
|-73.997177124023438|40.742168426513672|
|-74.006843566894531|40.730266571044922|
+-----+-----+
```

only showing top 5 rows

- DataFrames cast to RDDs return RDDs of rows

```
scala> val dist = df.select("trip_distance").rdd.map({ case Row(d: String) => d.toDouble }).mean
dist: Double = 15.7432129086462
```

Spark MLlib

- Spark's machine learning library
- Spark MLlib - RDD API (code maintenance only)
- Spark ML - DataFrame API
- Provides many common clustering and classification algorithms

Spark MLlib

- Spark's machine learning library
- Spark MLlib - RDD API (code maintenance only)
- Spark ML - DataFrame API
- Provides many common clustering and classification algorithms

```
val training: DataFrame
```

```
val lr = new LogisticRegression()  
  .setMaxIter(10)  
  .setRegParam(0.3)  
  .setElasticNetParam(0.8)
```

```
// Fit the model  
val lrModel = lr.fit(training)
```

```
// Print the coefficients and intercept for logistic regression  
println(s"Coefficients: ${lrModel.coefficients} Intercept: ${lrModel.intercept}")
```


Image Classification

[spark-sem-classify on GitHub](#)




Image Classification

```
// Assemble feature columns into a single column
val assembler = new VectorAssembler()
    .setInputCols(Array("bse", "ca", "si", "al", "mg",
        "o", "s", "k", "na", "fe"))
    .setOutputCol("features")

// Index labels, adding metadata to the label column.
val labelIndexer = new StringIndexer()
    .setInputCol("label")
    .setOutputCol("indexedLabel")
    .fit(data)

// Specify a RandomForest model.
val rf = new RandomForestClassifier()
    .setFeaturesCol("features")
    .setLabelCol("indexedLabel")
    .setNumTrees(30)

// Convert indexed labels back to original labels.
val labelConverter = new IndexToString()
    .setInputCol("prediction")
    .setOutputCol("predictionLabel")
    .setLabels(labelIndexer.labels)

// Chain indexers and forest in a Pipeline
val pipeline = new Pipeline()
    .setStages(Array(assembler, labelIndexer, rf, labelConverter))
```

Image Classification

```
// We now treat the Pipeline as an Estimator, wrapping it in a CrossValidator instance.
// This will allow us to jointly choose parameters for all Pipeline stages.
// A CrossValidator requires an Estimator, a set of Estimator ParamMaps, and an Evaluator.
// Note that the evaluator here is a BinaryClassificationEvaluator and its default metric
// is areaUnderROC.
val cv = new CrossValidator()
    .setEstimator(pipeline)
    .setEvaluator(evaluator)
    .setEstimatorParamMaps(paramGrid)
    .setNumFolds(3) // Use 3+ in practice

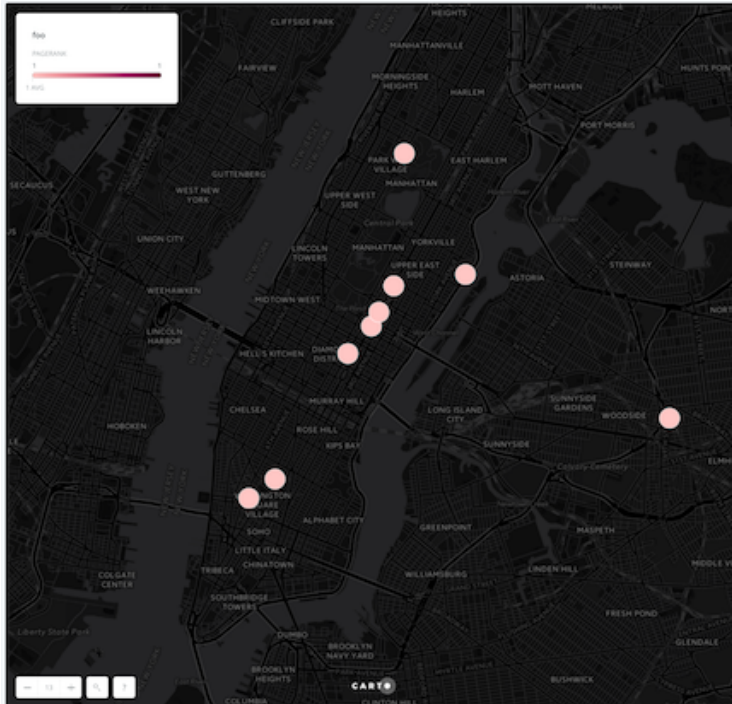
// Train model. This also runs the indexers.
val cvModel = cv.fit(trainData)
```

Spark GraphX

- Graph-parallel computation

Spark GraphX

- Graph-parallel computation
- For example, calculate the most popular taxi pickup/dropoff location in NYC using PageRank



Network Analysis

```
// Construct the edges
// Join the dropoff with it's index
val edges: RDD[Edge[String]] = pickups.map {
  case (_, (Row(_: Point, p1: Point), v0: Long)) =>
    (p1.round, v0)
} .join(indexed).values
  .map { case (v0: Long, v1: Long) => Edge(v0, v1, "foo") }

// Reformat indexed to be vertices
val vertices: RDD[(VertexId, Point)] = indexed.map({
  case (p: Point, v: Long) => (v, p)
})

// Specify default point
val defaultPoint = Point(-1.0, -1.0)

// Create the graph
val graph = Graph(vertices, edges, defaultPoint)

// Run PageRank
val ranks = graph.pageRank(0.0001) vertices

ranks take 10 foreach println
```

Pro Tips

1. Read the docs!

Pro Tips

1. Read the docs!
2. Find example code!!

Pro Tips

1. Read the docs!
2. Find example code!!
3. Use a good IDE!!!
 - type-checking
 - syntax highlighting
 - tab-completion
 - suggestions: [IntelliJ IDEA CE](#) and [PyCharm CE](#)